

Supplementary Material: When Agent Failure Prediction Measures Task Difficulty Rather Than the Run

ANONYMOUS AUTHOR(S)

This supplement collects derivations and full tables referenced in the main paper. Exhaustive per-cell grids, all seeds, and a make reproduce-main pipeline are in the replication package (<https://github.com/seanonymoussubmission/agent-monitor-estimand>); each item below names the script that produces it.

1 DERIVATIONS

1.1 The pooled decomposition (main §2.3)

Pooled AUROC is the concordance probability over positive–negative (failure–success) pairs. Partitioning pairs into same-unit and cross-unit gives the identity

$$A_{\text{pool}} = w A_{\text{within}} + (1 - w) A_{\text{cross}}, \quad w = \frac{\sum_u P_u N_u}{(\sum_u P_u)(\sum_u N_u)},$$

with P_u, N_u the failing/succeeding runs of unit u . It holds for any scoring function; we verify it to 10^{-9} in every reported configuration. For n units of r runs, same-unit pairs number $O(nr^2)$ against $O(n^2 r^2)$ cross-unit, so $w = \Theta(1/n)$; measured values are 2.5×10^{-4} (C1), 3.4×10^{-5} (C2), 8.8×10^{-5} (C3), 6.8×10^{-4} (C4-L).

1.2 The difficulty-oracle ceiling (main §2.3)

Score runs by the unit failure rate q_u . A drawn failure comes from unit a with weight $\propto r_a q_a$, a drawn success from b with weight $\propto r_b(1 - q_b)$. Writing $q_a(1 - q_b) = \frac{1}{2}[(q_a + q_b - 2q_a q_b) + (q_a - q_b)]$, the symmetric term integrates to $\mu(1 - \mu)/2$ over the ordered half and the antisymmetric term to $\text{MD}/4$ ($\text{MD} = \mathbb{E}|q_a - q_b|$, the Gini mean difference), giving

$$A_{\text{cross}}^{\text{oracle}} = \frac{1}{2} + \frac{\text{MD}}{4\mu(1 - \mu)}, \quad w = \frac{\overline{q(1 - q)}}{n\mu(1 - \mu)}.$$

Both are computable from the outcome table alone. Predicted vs. measured leave-one-run-out oracle: 0.954/0.945 (C1), 0.987/0.974 (C2), 0.984/0.903 (C3), 0.973/0.949 (C4-L), 0.979/0.959 (C4-Q). Script: `90_predict_inflation.py`.

1.3 Finite-run refinement (main §2.3)

Conditioning on a drawn run’s own label leaves $r - 1$ i.i.d. Bernoulli(q_u) runs, so both a drawn failure’s and a drawn success’s leave-one-out score are $\text{Bin}(r - 1, q_u)/(r - 1)$. The LOO oracle’s cross-unit AUROC is the concordance of the failure- vs. success-weighted mixtures of these; same-unit pairs are fully discordant (the anticorrelation artifact of the LOO estimator). This halves the ceiling’s error on all five corpora (0.950/0.980/0.940/0.962/0.969 predicted); the residual falls with runs per unit (0.037 at C3’s three, 0.005–0.013 at ten or more) and is plug-in dispersion bias. Script: `91_finite_r_theory.py`.

1.4 Analytic permutation-null SD (main §2.5)

Under H_0 the within-unit Mann–Whitney statistic has exact null variance $(P_u + N_u + 1)/(12P_u N_u)$; the pair-weighted aggregate has null SD $\sqrt{\sum_u P_u N_u (P_u + N_u + 1)/12 / \sum_u P_u N_u}$. Predicted vs. empirical:

Table 1. C4-L activation probe, within-task AUROC by prefix k and layer (142 mixed units). No cell survives Holm; best $k=10$, L30 (0.546, $p_{\text{Holm}} = 0.40$). Script: 24_lph_analysis.py.

	L10	L20	L30	L40
$k=1$	0.482	0.524	0.507	0.521
$k=3$	0.514	0.517	0.478	0.505
$k=5$	0.488	0.527	0.516	0.490
$k=10$	0.546	0.536	0.546	0.520

Table 2. C3 factorial: pre-execution predictors, held-out cross-fitted by task (within = 0.5 by construction). Script: 93_factorial_and_taskboot.py.

Predictor	pooled AUROC
model identity only	0.568
task only (difficulty $\times$ domain)	0.826
task $\times$ model	0.845
unit oracle (leave-one-run-out)	0.903
per-model task oracle (median; range)	0.985 (0.962–0.998)

Table 3. Operating curve, C4-L $k=5$ automaton monitor: abort when score $\geq \tau$. Script: 94_operational.py.

success retention	false-abort	failure recall	compute saved
0.99	0.01	0.014	0.012
0.97	0.03	0.054	0.043
0.95	0.05	0.091	0.073
0.90	0.10	0.139	0.120

0.0058 vs. 0.0057 (C2), 0.0183 vs. 0.0177–0.0190 (C3), 0.018 vs. 0.018 (C4-Q), 0.020 vs. 0.021 (C4-L probe subset). Script: 92_analytic_null.py.

1.5 Fluid-limit allocation anchor (main §6.2)

With known parameters the expected cost to solve u is $\bar{c}_u/\theta_u$; serving arms greedily by $\theta_u/\bar{c}_u$ maximises solves under budget in the fluid limit—a ratio-index policy (cf. Smith’s rule for weighted completion time). The oracle implements it and Thompson sampling approximates it in posterior form. Arm retirement on success is endogenous, outside the standard budgeted- and mortal-bandit analyses, so no regret guarantee is claimed.

1.6 Look-ahead leakage proposition (main §4.5)

If the cut index $k = g(T)$ is a function of eventual length T with g injective on the support, then observing the prefix reveals k and hence T ; any measurable function of T is a realisable score, so attainable within-task AUROC is lower-bounded by $\text{AUROC}(T) = 0.614$ (measured, C1). For a fixed index k , prefix statistics are functions of the first k turns only and carry no such bound.

2 FULL TABLES

Capability-legibility robustness (main, depth-and-capability analysis). The C3 capability trend ($\rho = -0.662$) survives every check we ran. Permutation: shuffling capability labels across the

Table 4. Label-noise calibration (C4-L): observed within-task AUROC after corrupting a fraction α of success labels, for synthetic signals of three true effect sizes. Script: 81_noise_calibrated.py.

α	true 0.55	true 0.60	true 0.65
0.00	0.542	0.614	0.664
0.05	0.529	0.577	0.620
0.10	0.521	0.560	0.600
0.20	0.514	0.542	0.565
0.30	0.507	0.528	0.547
0.50	0.508	0.524	0.539

Table 5. Bandit allocation, C4-Q target (prior transferred from C4-L), distinct tasks solved by token budget; 200 trials. Script: 54_bandit_allocation.py.

Policy	2M	5M	10M	40M
Round-robin	53.1	131.3	261.7	360.3
Static ranking, transferred prior	157.2	208.0	208.0	—
Bandit, uninformative prior	135.4	212.7	269.9	358.2
Bandit + transferred prior	162.5	246.0	307.9	360.5
Oracle (true θ_u)	172.3	277.2	345.6	375.0

Table 6. Cold-start prior seeded from pre-execution features, novel-task regime (C4-L), tasks solved at scaled budget vs. pseudo-count weight. Script: 68_coldstart_prior.py.

prior	tasks solved
uninformative	108.99 $\pm$ 4.45
pre-execution seed, $w=0.25$	108.63
$w=1.0$	107.35
$w=4.0$	102.36
$w=8.0$	96.90
transferred (rollout-derived)	124.74
oracle	139.10

Table 7. C1 depth sweep on a fixed cohort of 63,657 traces (885 mixed units, identical at every k), automaton featuriser (main, depth-and-capability analysis). The varying-sample column is what a naive sweep reports; the two agree, so selection is not driving the rise. Script: 66_characterize.py.

k	varying sample	fixed cohort (95% CI)
1	0.5072	0.5088 [0.4982, 0.5189]
3	0.5114	0.5105 [0.4984, 0.5211]
5	0.5386	0.5320 [0.5214, 0.5428]
10	0.5858	0.5858 [0.5717, 0.5988]

17 models gives a null centred at +0.0002 (sd 0.247), exact $p = 0.0046$. Confounds: partialling out outcome-blind behavioural dispersion moves the estimate only $-0.662 \rightarrow -0.668$. Dependence: a model-resample bootstrap that accounts for tasks shared across models gives a 95% CI of

Table 8. Within-unit concordance by capability tercile on C3, each behavioural family fitted alone (main, depth-and-capability analysis). Weak models are legible through several channels; for stronger models none clears chance. Script: 73_signal_source_by_capability.py.

Family	weak	middle	strong
All families	0.666	0.521	0.529
Transition surprise	0.638	0.508	0.483
Output size	0.610	0.504	0.484
Which actions	0.609	0.534	0.487
Repetition	0.534	0.474	0.477
Errors seen	0.529	0.470	0.501

Table 9. Abort-and-restart deployment simulation on C1 (main, abort-vs-reallocate analysis): a restart redraws the outcome from the same exchangeable unit; baseline failure rate 0.8334. Δ success, disruption/recovery ratio d/r , and break-even range span all operating points. The predictor selected by A_{pool} is negative at all seven points swept; the one selected by A_{within} is non-negative at all twelve. Script: 79_abort_vs_allocate.py.

Predictor (selection metric)	Δ success	d/r	break-even
Oracle (A_{pool})	−0.0012 to −0.0028	14.7–348	0.936–1.000
Trajectory (A_{within})	+0.0000 to +0.0029	5.1–21.8	0.835–0.956

[−0.90, −0.23], and a precision-weighted slope is -0.674 ± 0.113 . Measurement: capability computed on disjoint tasks and on disjoint domains, matched within-unit pair counts, and leave-one-family-out all preserve the negative sign. Script: 72_behavioral_dispersion.py, 74_within_family.py, 85_dependence.py.

3 VERIFICATION PROTOCOL

Every probe result passes four gates (main §5.8): (1) weight integrity—zero missing keys and a teacher-forced accuracy floor of 0.55 (one code path silently randomised all 39 MoE layers; Apple-MPS loaded correctly yet computed wrong values); (2) replay from stored token ids, never a reconstructed template; (3) an independent estimator implementation agreeing to 10^{-9} ; (4) positive controls—an outcome-carrying feature recovers within-task 1.000, Gaussian noise 0.503, and real activations with shuffled labels 0.488. Scripts: 27_load_integrity.py, 61_validate.py.

4 LITERATURE-SEARCH PROTOCOL

Four Google Scholar query families (failure attribution; LLM-agent failure prediction; agent-trajectory failure; coding-agent failure), 2023 onward; every returned trajectory-level failure-prediction result was read and snowballed, and each load-bearing claim verified against the cited paper’s own text. The citation-knee stopping rule is inapplicable to a literature whose 2026 works have zero citations, so we substituted an exhaustive read of the returned set.